

Exercise HandsOn – Stream Processing

Advanced DBMS

MIT 2019 - Data Science

June 2020



Dennis Timis
Bojan Makivic
Johannes Wieser

Team	3
Members	3
Work distribution	3
Description of Use Case	4
Web-APIs for data retrieval	5
Twitter	5
Yahoo! Finance	5
Overview of the solution	6
High Level Architecture Overview	6
Used Technologies	6
Python	6
Kafka	6
Installation & Setup	9
Twitter Developer Accounts	9
Kafka Setup	10
Messages	14
Python Environment	15
How to run the project	15
Python libraries	16
confluent_kafka	16
twitter	16
yfinance	16
json	16
matplotlib	16
nltk (Natural Language Toolkit)	17
Task 1 - Producer	18
twitter_api.py	18
yahoo_api.py	18
requestor.py	19
producer.py	19
producer_factory.py	19
main.py	20
Task 2 - Consumer	21
Connection	21
Task 3 - Visualization	24

1. Team

1.1. Members

Name	Student #
Bojan Makivic	01002146
Dennis Timis	51818530
Johannes Wieser	00608831

1.2. Work distribution

Task	Team Member
Task 1	Dennis Timis
Task 2	Johannes Wieser
Task 3	Bojan Makivic

2. Description of Use Case

Our Idea to fulfil the requirement of using several data sources to combine and stream them with a stream processor was to use Tweets and financial Stock data and combine them to find out if there are some correlations. So the specific example looks like this:

- Continuously get stock data on Microsoft Stock
- Continuously get stock data on Tesla Stock
- Continuously get all tweets containing the word "Gates"
- Continuously get all tweets containing the word "Musk"

As we wanted to find out a relation between twitter and stock market price we took the following metrics:

- Stock market price in USD
- Amount of tweets per time
- Calculate average sentiment of the tweets per time (from -1 to 1)

Later (see [visualization](#)) we combined and compared the output while it's dynamically updating.

3. Web-APIs for data retrieval

3.1. Twitter

In order to use the Twitter API, first a developer account needs to be set up. This procedure will be described in [chapter 5.1](#). However, once everything is properly set up, you can easily engage with the Twitter API and search tweets, post content and so on.

Twitter offers three different API modalities:

- Standard Search API
- Premium Search API
- Enterprise Search API

For this project the Standard Search API was taken because it is useable for free (check the restrictions, though!) and sufficient for the use cases.

For a more precise documentation refer to the documentation:

<https://developer.twitter.com/en/docs/tweets/search/api-reference/get-search-tweets>

In order to avoid doing the HTTP Post and Receive action by ourselves, we used an existing twitter library for Python called `twitter`. See [chapter 5.3.2](#) for more information.

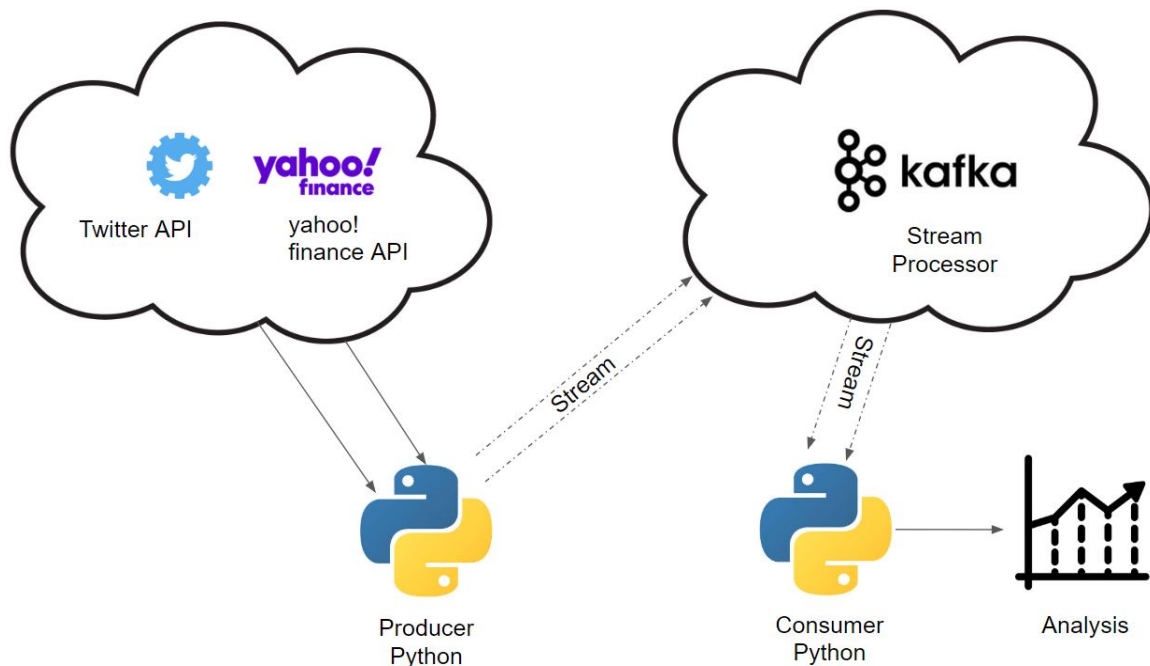
3.2. Yahoo! Finance

For the retrieval of the stock prices Yahoo! Finance was used because it is not only simple to use, but it also does not need any user account, authentication or whatsoever. Basically all you need to do is perform your request and you will get your result.

For this the library `yfinance` was used. See [chapter 5.3.3](#) for more information.

4. Overview of the solution

4.1. High Level Architecture Overview



The overall approach high level overview consists of all systems involved in the project. The Data is collected from the twitter and yahoo finance APIs and distributed via python in the stream processing system, which is hosted in the cloud. The data is then streamed continuously and later collected and polled with python and analyzed.

4.2. Used Technologies

4.2.1. Python

Python was used to get the data from the Twitter and Yahoo finance APIs and produce Messages to the kafka stream processor. It was also used to poll the data out of the kafka in the form of a stream and create upon this analysis, in kafka this is called the “consumer”.

4.2.2. Kafka

Apache Kafka is an open-source stream-processing software platform developed by the Apache Software Foundation, written in Scala and Java. The project aims to provide a unified, high-throughput, low-latency platform for handling real-time data feeds. Its storage layer is essentially a “massively scalable pub/sub message queue architected as a distributed transaction log,” making it highly valuable for enterprise infrastructures to process

streaming data. Additionally, Kafka connects to external systems (for data import/export) via Kafka Connect and provides Kafka Streams, a Java stream processing library. (Wikipedia)

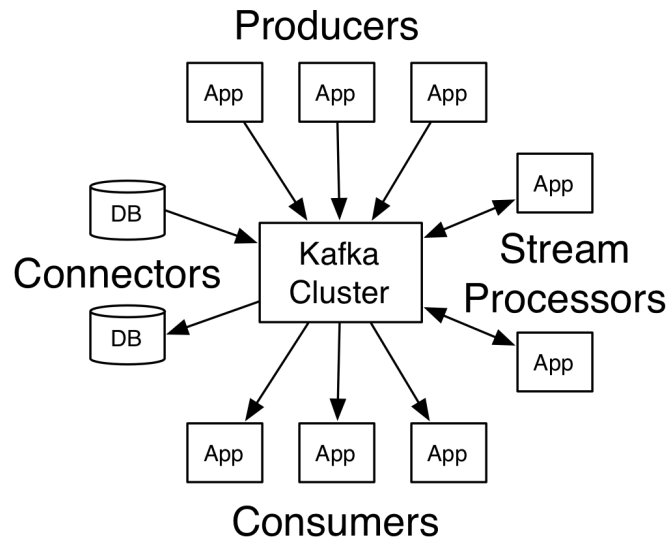


Image Source: Kafka

Topics

Every message that is fed into the system must be part of some topic. The topic is nothing but a stream of records. The messages are stored in key-value format. Each message is assigned a sequence, called Offset. The output of one message could be an input of the other for further processing.

Producers

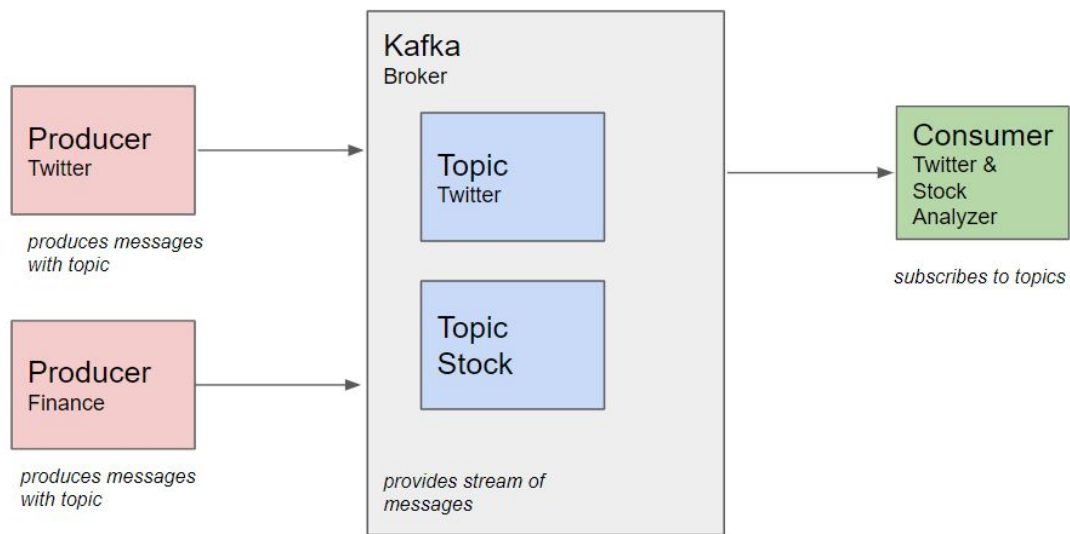
Producers are the apps responsible to publish data into the Kafka system. They publish data on the topic of their choice.

Consumers

The messages published into topics are then utilized by Consumers apps. A consumer gets subscribed to the topic of its choice and consumes data.

Broker

Every instance of Kafka that is responsible for message exchange is called a Broker. Kafka can be used as a stand-alone machine or a part of a cluster.



In our case, the Producers create Messages for Stock Data, as well as Tweets from Twitter and feed them into Kafka which is hosted in the cloud using Confluent Kafka. The Consumer can choose which Topics to subscribe to, in our case both Topics "Twitter" and "Stock" are used in the consumer application. In the Consumer the data is analyzed and visualized dynamically as the data changes continuously.

5. Installation & Setup

5.1. Twitter Developer Accounts

The procedure how to set up a developer account is pretty straight forward and explained here: <https://developer.twitter.com/en/docs/basics/developer-portal/overview>

Once you have the developer account, you need to know your login credentials in order to interact with the API. For this follow these steps:

- 1) Login and access <https://developer.twitter.com/en/apps> .
- 2) Choose your app and click on “Details”..
- 3) Navigate to the “Keys and token” tab.

Here we are presented with the “Consumer API keys” which are already existing. What needs to be done is the generation of the “Access token & access token secret”.

- 4) Click on the “Generate” button in the section “Access token & access token secret”.

Once you have generated them, store them somewhere secretly (if needed) because they will be hidden afterwards! With these 4 tokens you can now use the twitter API.

The screenshot shows the 'Keys and tokens' page in the Twitter Developer Portal. At the top, there are tabs for 'App details', 'Keys and tokens' (which is selected), and 'Permissions'. The main content area is titled 'Keys and tokens' with a subtitle 'Keys, secret keys and access tokens management.' Below this, there are two main sections. The first section is 'Consumer API keys', which shows an 'API key' and an 'API secret key' (both masked with dots). There is a 'Regenerate' button to the right. The second section is 'Access token & access token secret', which shows an 'Access token', an 'Access token secret' (both masked with dots), and an 'Access level' of 'Read and write'. There are 'Revoke' and 'Regenerate' buttons to the right. A note below this section states: 'We only show your access token and secret when you first generate it in order to make your account more secure. You can revoke or regenerate them at any time, which will invalidate your existing tokens.' The 'Access token' and 'Access token secret' are both masked with dots, and the 'Access level' is 'Read and write'. The 'Last generated' date is 'Jun 4, 2020'.

5.2. Kafka Setup



Apache Kafka® on Confluent Cloud™

Confluent

Apache Kafka Re-Engineered for the Cloud

ENABLE ⓘ [Purchased on 6/4/20 but not yet enabled](#)

OVERVIEW PRICING SUPPORT

On the Google Cloud Platform (GCP), Apache Kafka was enabled and can be billed with your Google Cloud account. We registered a collectively used gmail account to use GCP services.

After Registration a new Kafka Cluster is created:

New cluster

Cluster name* ⓘ
cluster_0

Provider & region ⓘ

Google Cloud Platform

Region*
europe-west3

[Request a region](#)

Cluster type ⓘ

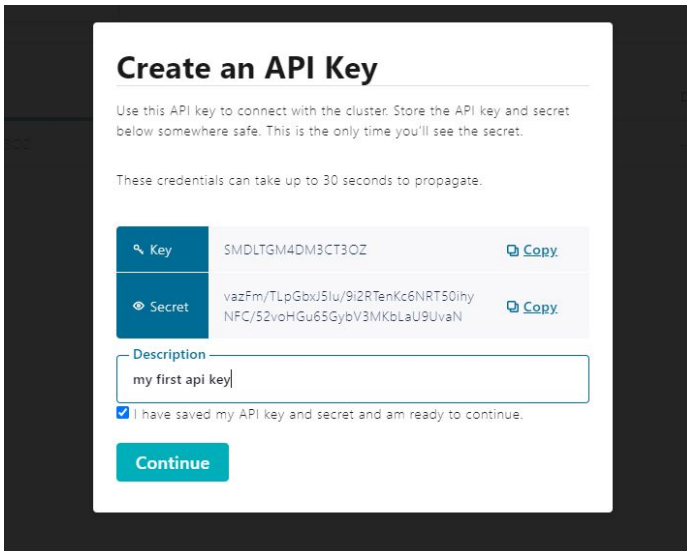
Basic Standard

Availability ⓘ

Multi zone Single zone

Continue Cancel

And the cluster is later launched. After this the API Key and secret were generated. These are required for external access to the Kafa Cloud Cluster.



The screenshot shows a 'Create an API Key' form. At the top, it says 'Create an API Key'. Below that, a message states: 'Use this API key to connect with the cluster. Store the API key and secret below somewhere safe. This is the only time you'll see the secret.' Another message says: 'These credentials can take up to 30 seconds to propagate.' The form contains two rows: 'Key' with the value 'SMDLTGM4DM3CT3OZ' and a 'Copy' button, and 'Secret' with the value 'vazFm/TLpGbxJ5lu/9l2RTenKc6NRT50lhy/NfC/\$2voHG65GybV3MKbLaU9UvaN' and a 'Copy' button. There is a 'Description' field with the text 'my first api key'. A checkbox is checked with the text 'I have saved my API key and secret and am ready to continue.' At the bottom is a 'Continue' button.

Type	Value	Action
Key	SMDLTGM4DM3CT3OZ	Copy
Secret	vazFm/TLpGbxJ5lu/9l2RTenKc6NRT50lhy/NfC/\$2voHG65GybV3MKbLaU9UvaN	Copy

Description: my first api key

☒ I have saved my API key and secret and am ready to continue.

Continue

You can log in and monitor the system on <https://confluent.cloud/login>

Kafka Dashboard

The overall “check” for your Kafka cluster can be done on the “Overview” after the login. It gives instant information if data is streaming, which topics are in use, and basically lets you see if everything is working or if there are some Problems.



Topics

2

Total topics

Partitions[®]

6

Total partitions

Connectors

0

Total

ksqlDB Applications

0

Total

Message View per Topic

In real time the messages per topic can be viewed, and also messages can be produced for debugging etc. It's accessed via Topics > "topic" > Message

The screenshot shows the Kafka Message View per Topic interface. At the top, there is a search bar labeled "Filter by keyword" and a "Jump to offset" dropdown menu. Below the search bar, there is a button labeled "Produce a new message to this topic". The main area displays a list of messages, each with a timestamp, tweet ID, username, and partition/offset information. The messages are sorted by timestamp, with the most recent message at the top. A "Newest" button is visible on the right side of the message list.

Timestamp	Tweet ID	Username	Partition	Offset	Timestamp
2020-06-18 13:57:36	1273721512885878800	annakarienin	1	70573	1592513856497
2020-06-18 13:57:36	1273721525187862500	WienerKoks	1	70572	1592513856497
2020-06-18 13:57:36	1273721526399860700	RaeJ03	1	70571	1592513856497
2020-06-18 13:57:36	1273721531697434600	Mimchen	1	70570	1592513856497
2020-06-18 13:57:36	1273721533391876000	justicewarriorD	1	70569	1592513856497

A very useful feature is also that a message can be produced here.

Streaming Window

Per Topic the streaming window size can be configured. Either in maximum amount of bytes or in time. We choose to leave all messages for 10 seconds in the window. The configuration can be found in Confluent Kafka > Topics > Configuration > Switch to expert mode > retention_ms was set to 10000. For a retention maximum byte size, "retention_bytes" can be used. We set it to "-1" which is "maximum"

The screenshot shows the Kafka Configuration page for a topic. The configuration is displayed in a table with the following values:

Configuration Name	Value
confluent_value_subject_name_strategy	io.confluent.kafka.serializers.subject.TopicNameStrategy
retention_ms	10000
segment_bytes	104857600
flush_messages	9223372036854775807
message_format_version	2.3-IV1

5.3. Messages

5.3.1. Overall

The messages are in json format and handled via Kafka. We have agreed on, but not enforced a schema on each topic's messages.

5.3.2. Stock Messages

The stock messages contain these values:

- timestamp: the timestamp created by the producer when the data is received
- value: stock price in USD
- stock_timestamp: timestamp from the yahoo! finance api

Example Message

```
{
  "timestamp": "2020-06-19 07:49:31",
  "value": 198.12,
  "stock_timestamp": "2020-06-19 10:49:20-04:00"
}
```

5.3.3. Twitter Messages

The twitter messages contain these values:

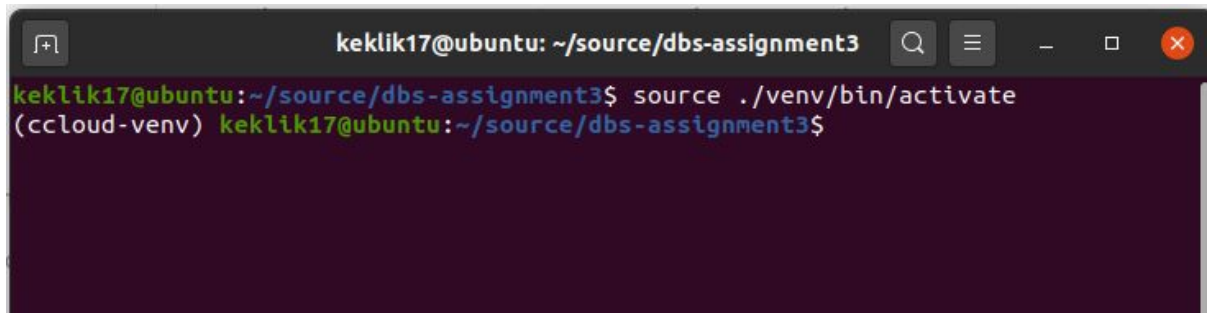
- timestamp: the timestamp created by the producer when the data is received
- tweet
 - id: tweet unique id
 - username: twitter username from the tweet author
 - tweet_timestamp: twitters timestamp
 - text: the content of the tweet

Example Message

```
{
  "timestamp": "2020-06-19 07:51:00",
  "tweet": {
    "id": 1273991666022731800,
    "username": "BornintheUSA100",
    "tweet_timestamp": "Fri Jun 19 14:50:57 +0000 2020",
    "text": "RT @MzMugzzi: This is the real Gates! https://t.co/gmsZVgsp36"
  }
}
```

5.4. Python Environment

A virtual Environment was used to install all a version of python and all required libraries to run this project on your machine. It is not included in the project because it's too big for edunet upload, but can be downloaded [here](#). It can be activated with the activation script. Required for the use is only "python", "pip" and "virtualenv". All other packages as well as the specific are included in the venv folder, and used from the terminal in which it's activated. See the name of the environment in brackets before the command line:

A terminal window with a dark background. The title bar shows 'keklik17@ubuntu: ~/source/dbs-assignment3'. The prompt is 'keklik17@ubuntu:~/source/dbs-assignment3\$'. The command 'source ./venv/bin/activate' has been entered. The output shows the prompt changed to '(ccloud-venv) keklik17@ubuntu:~/source/dbs-assignment3\$'.

5.5. How to run the project

Configure the in the main.py which producers you want to choose.

```
factory.register_twitter('Musk')
factory.register_twitter('Gates')
factory.register_yahoo('TSLA')
factory.register_yahoo('MSFT')
```

- Install python, pip and virtualenv
- Download the virtual environment with all pre-installed packages [here](#)
- Open a terminal
- Activate the virtual environment by using "source path/venv/bin/activate"
- run "python path/main.py"
- you see the producers starting to fire messages to kafka
- Open a new terminal
- Activate the virtual environment by using "source path/venv/bin/activate"
- run "python path/twitter-analysis-consumer.py"

5.6. Python libraries

5.6.1. confluent_kafka

This python library is a client which abstracts the communication with the Kafka server. The documentation on how to work with the client is found here:

<https://docs.confluent.io/current/clients/confluent-kafka-python/>

In order to install this library on your machine open the command-line and run:

```
pip install confluent-kafka
```

5.6.2. twitter

This python library is a wrapper around the Twitter API. The documentation on how to work with this wrapper is found here: <https://python-twitter.readthedocs.io/en/latest/>

In order to install this library on your machine open the command-line and run:

```
pip install python-twitter
```

5.6.3. yfinance

This python library is a wrapper around Yahoo's stock market API. The documentation on how to work with this wrapper is found here: <https://pypi.org/project/yfinance/>

In order to install this library on your machine open the command-line and run:

```
pip install yfinance
```

5.6.4. json

This python library is responsible for the conversion of Python data to a json string (that is the format the Kafka server reliably works with) and vice versa. The documentation on how to work with the client is found here: <https://docs.python.org/3/library/json.html>

This library does not need to be separately installed because it is part of the standard library.

5.6.5. matplotlib

This python library is a comprehensive library for creating static, animated, and interactive visualizations. The documentation on how to work with this library is found here:

<https://matplotlib.org/contents.html>

Regarding the installation it depends on your exact environment. Take look at this documentation: <https://matplotlib.org/users/installing.html#building-on-windows>

5.6.6. nltk (Natural Language Toolkit)

This python library is a leading platform for building Python programs to work with human language data. One of its relevant functionalities for this project is the ability to perform sentiment analysis. The documentation on how to work with this library is found here:

<https://www.nltk.org/>

In order to install this library on your machine open the command-line and run:

```
pip install nltk
```

6. Task 1 - Producer

This task represents the creation of a producer, using the `confluent_kafka` client-library, to periodically send content to the Kafka server. The producer is split into multiple projects files which will be not shortly explained.

6.1. twitter_api.py

This module holds the functionality to interact with the twitter API.

```
class TwitterAPI(Requestor):

    # constructor
    # @ authentication: Represents the Twitter credentials.
    # @ search_word: Represents the name after which to be searched.
    def __init__(authentication: TwitterAuthentication,
                  search_word: str) -> TwitterAPI

    # Implementation of the abstract method to retrieve new data.
    # - returns: Returns a List of dictionaries containing the search result.
    def request_new() -> [{}]
```

```
class TwitterAuthentication():

    # constructor
    # @ consumer_key: Represents Twitter's "Consumer Key".
    # @ consumer_secret: Represents Twitter's "Consumer Secret".
    # @ access_token_key: Represents Twitter's "Access Token Key".
    # @ access_token_secret: Represents Twitter's "Access Token Secret".
    def __init__(consumer_key: str, consumer_secret: str, access_token_key: str,
                  access_token_secret: str) -> TwitterAuthentication
```

6.2. yahoo_api.py

This module holds the functionality to interact with the yahoo API.

```
class YahooFinanceAPI(Requestor):

    # constructor
    # @ stock: Represents the name of the stock to be analyzed.
    def __init__(stock: str) -> YahooFinanceAPI

    # Implementation of the abstract method to retrieve new data.
    # - returns: Returns a List of dictionaries containing the search result.
    def request_new() -> [{}]
```

6.3. requestor.py

This module holds the abstract class which needs to be implemented by the twitter API and the yahoo API.

```
class Requestor():  
  
    # Represents the abstract method which needs to be implemented by all subclasses.  
    # - returns: Returns a List of dictionaries containing the search result.  
    @abstractmethod  
    def request_new() -> [{}]
```

6.4. producer.py

This module wraps a requestor and periodically calls its interface method to retrieve new data. The class is run in a separate thread, so that the main thread is not blocked.

```
class Producer():  
  
    # constructor  
    # @ host: Represents the host name of the Kafka server.  
    # @ port: Represents the port where the Kafka server is run.  
    # @ username: Represents the username for authentication.  
    # @ password: Represents the password for authentication.  
    def __init__(host: str, port: int, username: str, password: str) -> Producer  
  
    # Starts the producer to periodically retrieve new data.  
    def start(requestor: Requestor,  
              topic_name: str,  
              topic_key: str,  
              timeout: int) -> [{}]  
  
    # Stops the producer.  
    def stop() -> None
```

6.5. producer_factory.py

This module represents a factory, holding and enabling the creation of multiple producers. First of all the producers need to be created and afterwards the factory can start the registered producers. Once the factory has initiated the start, no further producers can be registered to be run. Stopping the factory will reset it and remove all registered producers.

```
class ProducerFactory():  
  
    # constructor  
    # @ host: Represents the host name of the Kafka server.
```

```

# @ port: Represents the port where the Kafka server is run.
# @ username: Represents the username for authentication.
# @ password: Represents the password for authentication.
def __init__(host: str, port: int, username: str, password: str)
    -> ProducerFactory

# Starts all the registered producers.
def start() -> None

# Stops the the registered producers and empties the registered List.
def stop() -> None

# Registers a producer to retrieve Twitter content.
# @ search_word: Represents the name after which to be searched.
def register_twitter(search_word: str, ) -> None

# Registers a producer to retrieve Twitter content.
# @ stock: Represents the name of the stock to be analyzed.
def register_yahoo(stock: str, ) -> None

```

6.6. main.py

This module represents the entry point of the producer's application and contains the credentials for the authentication on the Kafka server.

```

import producer
import producer_factory
import requestor
import twitter_api
import yahoo_api

# Represents the credentials for the Kafka server.
HOST: str = '***.gcp.confluent.cloud'
PORT: int = 9092
USER_TOKEN: str = '***your password token here***'
PASSWORD_TOKEN: str = '***your password token here***'

# Created the factory.
factory = producer_factory.ProducerFactory(HOST, PORT, USER_TOKEN, PASSWORD_TOKEN)

# Registers producers for twitter and yahoo! finance.
factory.register_twitter('Tesla')
factory.register_twitter('Gates')
factory.register_yahoo('TSLA')
factory.register_yahoo('MSFT')

# Starts the factory.
factory.start()

```

7. Task 2 - Consumer

The purpose of the consumer is to get all messages which have the required topics - which in this case is “stock” and “twitter”. And further prepare and hold the data to later visualize it.

7.1. Connection

The connection is established using the `confluent_kafka` package and the key and secret as well as configuration data.

```
from confluent_kafka import Consumer

# create consumer and establish connection
c = Consumer({
    'bootstrap.servers': '***.confluent.cloud:9092',
    'sasl.mechanism': 'PLAIN',
    'security.protocol': 'SASL_SSL',
    'sasl.username': '***username***',
    'sasl.password': '***password***',
    'group.id': str(uuid.uuid1()), # this will create a new consumer group on each
    invocation.
    'auto.offset.reset': 'earliest'
})
```

7.2. Subscription to Topics & Polling

The package also makes it very simple to subscribe to one or several topics it's.

```
# subscribe to both topics
c.subscribe(['twitter', 'stock'])
```

After the subscription it's easily possible to poll messages. We do this in a while True loop to continuously poll for data.

```
try:
    while True:
        msg = c.poll(0.1) # Wait for 0.1 secs for message
        if msg is None:
            # No message available within timeout.
            # Initial message consumption may take up to `session.timeout.ms` for
            # the group to rebalance and start consuming.
            continue
        if msg.error():
            # Errors are typically temporary, print error and continue.
            print("Consumer error: {}".format(msg.error()))
            continue
```

It is also continuously checked for errors or null messages.

7.3. Message Selection

```
if msg.topic() == 'twitter':
    if str(msg.key(), 'utf-8').lower() == 'tesla':
        myjson = json.loads(str(msg.value(), 'utf-8'))
        list_tweets_tsla.append(myjson['tweet']['text'])
        count_tweets_tsla +=1

    if str(msg.key(), 'utf-8').lower() == 'gates':
        myjson = json.loads(str(msg.value(), 'utf-8'))
        list_tweets_ms.append(myjson['tweet']['text'])
        count_tweets_ms +=1
```

Basically the messages are selected by Key and Topic and separated and temporarily saved to lists, which are later used for the visualization.

7.4. Sentiment Analysis

On all tweets a sentiment analysis is performed using nltk. To get the pretrained model for that on the first time a lexicon is downloaded.

```
import nltk
nltk.download('vader_lexicon')
from nltk.sentiment.vader import SentimentIntensityAnalyzer

#sentiment analyzer
sid = SentimentIntensityAnalyzer()
```

After that the sentiment is easily calculated by using `sid.polarity_scores("some text")['property']`. There are various properties. for our use the compound score gives us the most easily usable result which can vary from -1 (very bad sentiment) to 1 (very good sentiment). Every tweet text is analyzed, and the sentiment is averaged across all tweets.

```
if str(msg.key(), 'utf-8').lower() == stock_ms.lower():
    list_tweet_counts_ms.append(count_tweets_ms)
    count_tweets_ms = 0
    myjson = json.loads(str(msg.value(), 'utf-8'))
    list_stock_prices_ms.append(myjson['value'])
```

```

if len(list_tweets_ms) > 0:
    temp = []
    for tweet in list_tweets_ms:
        temp.append(sid.polarity_scores(tweet)['compound'])
    list_sentiments_ms.append(sum(temp) / len(temp))
    list_tweets_ms = []
else:
    list_sentiments_ms.append(0)

```

Also the sentiments are continuously hold in memory in lists to later be used in visualization to analyze if the tweets for the words are “good” or “bad”.

7.5. Tweet counting & Stock Price

Similarly but less complicated then the sentiment, the stock price and count of tweets are aggregated and saved in memory in lists. these lists are later used for visualization.

```


if msg.topic() == 'twitter':

if str(msg.key(), 'utf-8').lower() == twitter_searchword_tsla.lower():
    myjson = json.loads(str(msg.value(), 'utf-8'))
    list_tweets_tsla.append(myjson['tweet']['text'])
    count_tweets_tsla +=1

```

8. Task 3 - Visualization

We have used the matplotlib package for further analysis and visualization of streaming data. The aim was to visualize four graphs; two graphs with different stock market prices for two different acteurs and two graphs considering the count of keywords on twitter that we have previously defined. In addition the sentiment values were integrated in the two lower graphs filled the three categories by three different colors (figure below).



The biggest challenge was the implementation of colored sentiment values in the curved area due to non-user friendly possibilities that matplotlib offers for such visualization. The one possibility was to scale the 'y axis' values from 0-1 and to add 0.03 more points in order to fill fully the area under the curve as in the code snippet below.

```
axs[1,0].fill_between(range(1,len(list_sentiments_ms)+1),[i/max(np_tweet_counts_ms)+0.03
for i in np_tweet_counts_ms],where=np_sent_ms > threshold, color='xkcd:green',
alpha=0.9, transform=axs[1, 0].get_xaxis_transform())

axs[1,0].fill_between(range(1,len(list_sentiments_ms)+1),
[i/max(np_tweet_counts_ms)+0.03 for i in np_tweet_counts_ms], where=np_sent_ms <
threshold, color='xkcd:red', alpha=0.9, transform=axs[1, 0].get_xaxis_transform())
```



```
axs[1,0].fill_between(range(1,len(list_sentiments_ms)+1),[i/max(np_tweet_counts_ms)+0.03
for i in np_tweet_counts_ms], where=np_sent_ms == threshold, color='xkcd:blue',
alpha=0.9, transform=axs[1, 0].get_xaxis_transform())
```